

```
In [38]: %matplotlib inline
```

Underfitting vs. Overfitting

Model used: Linear regression with polynomial features to approximate nonlinear functions.

```
In [39]: import numpy as np # For numerical calculation and matrix handling
import matplotlib.pyplot as plt # For plotting

from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures # For pre-processing

from sklearn.linear_model import LinearRegression # For Linear regression model
from sklearn.model_selection import cross_val_score # For evaluation

np.random.seed(0) # To control the random number generator
```

Generate data for regression

```
In [40]: def gen_target(X):
return np.cos(1.5 * np.pi * X)
```

Define constants

```
In [41]: n_records = 5 # Total number of records
degrees = [1, 4, 30] # Degree(s) of linear regression model
types = ['Underfitting', 'Perfect fitting', 'Overfitting']
```

Generate features and targets

```
In [42]: X = np.sort(np.random.rand(n_records)) # Randomly generate data points (features)
y = gen_target(X) + np.random.randn(n_records) * 0.1 # Generate regression output
```

Build and Evaluate model

Plotting function

The plot shows the true function, the function approximated using linear regression model, and the records with additive noise used for building linear regression model. The model uses polynomial features of different degrees.

```
In [43]: def plot_test(X, y, deg, title=""):
X_test = np.linspace(0, 1, 100)
plt.plot(X_test, pipeline.predict(X_test[:, np.newaxis]), label="LR function (d
plt.plot(X_test, gen_target(X_test), '--r', label="True function")
plt.scatter(X, y, facecolor="b", s=20, label="Training records")
plt.xlabel("x")
plt.ylabel("y")
plt.xlim((0, 1))
plt.ylim((-2, 2))
plt.legend(loc="best")
plt.title(title)
```

- **Underfitting:** A linear function (polynomial with degree 1) is not sufficient to fit the training records.
- **Overfitting:** A polynomial of degree > 1 that approximates the true function almost perfectly and for higher degrees the model learns the noise of the training data.

Underfitting/Overfitting can be quantitatively evaluated by using cross-validation to calculate the mean squared error (MSE) on the validation set. The higher MSE, the less likely the model generalizes correctly from the training data.

```
In [44]: plt.figure(figsize=(14, 5)) # Generate figure window
for i, (deg, t) in enumerate(zip(degrees, types)):
    ax = plt.subplot(1, len(degrees), i + 1) # Generate subplot for each degree

    poly_feat = PolynomialFeatures(degree=degrees[i], include_bias=False)
    lr = LinearRegression()

    # Make regression pipeline
    pipeline = Pipeline(
        [
            ("poly_feat", poly_feat),
            ("lr", lr),
        ]
    )
    pipeline.fit(X[:, np.newaxis], y)

    # Evaluate the models using 10-fold cross-validation and MSE
    scores = cross_val_score(pipeline, X[:, np.newaxis], y, scoring="neg_mean_squar

    # Plot results with original data
    plot_test(X, y, deg)
    print("Degree {} \nMSE = {:.3e} (+/- {:.3e}) \n".format(deg, -scores.mean(), scor

plt.show()
```

Degree 1

MSE = $4.077\text{e-}01$ (+/- $4.255\text{e-}01$)

Degree 4

MSE = $4.321\text{e-}02$ (+/- $7.078\text{e-}02$)

Degree 30

MSE = $3.834\text{e+}10$ (+/- $1.059\text{e+}11$)

